

How legible is our community?:

LGBTQ+ Uses of LLM Chatbots in Mental Health, Sexuality, and Relationships

ANONYMOUS AUTHOR(S)

Large Language Models (LLMs) are widely used for personal care, from mental health guidance to emotional support. However, LLMs do not serve all communities equitably, often reproducing social bias when prompts carry demographic cues. Although bias evaluations typically rely on explicit identity markers, we hypothesize that LGBTQ+ users rarely disclose identity directly, signaling it instead through implicit or coded language. We investigate how distinguishable LGBTQ+ prompts are and what usage patterns and attitudes characterize this community's engagement with LLM chatbots. We survey LGBTQ+ and matched cisgender heterosexual respondents, eliciting ecologically valid prompts across mental health, gender and sexuality, and relationships, plus donated transcripts. We contribute an annotated cross-population dataset of naturalistic prompts (COMMUNITY-SUBTEXT); an analysis of the features that distinguish LGBTQ+ prompts, including lexical, syntactic, and semantic characterizations; and domain-level measures of usage frequency, aversion, and satisfaction. By centering the real prompting patterns of the LGBTQ+ community, our work informs the design of LLM chatbots that mitigate social bias across the full spectrum of how identity is expressed. As AI mediates more of daily life, how we are read from our prompts shapes how we are treated, making legibility a central concern for HCI.

ACM Reference Format:

Anonymous Author(s). 2026. How legible is our community?: LGBTQ+ Uses of LLM Chatbots in Mental Health, Sexuality, and Relationships . 1, 1 (July 2026), 19 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Therapy and companionship have remained the top uses of Large Language Models (LLMs) throughout 2025 and 2026 [69]. For minoritized populations, including the LGBTQ+ community, LLM chatbots widen access to mental health support [33, 35]. Yet these systems reproduce human biases [42, 49] and can ultimately harm the users they aim to help [17]. LLMs frequently underperform precisely when interacting with the vulnerable communities most in need of support [44, 46], and their safety mechanisms falter when identity is conveyed implicitly rather than explicitly [23, 25].

In this paper, we seek to answer the following questions:

- **RQ1:** How distinguishable are the prompts from members of LGBTQ+ community from cisgender heterosexual people?
- **RQ2:** What are the usage patterns and attitudes from members of LGBTQ+ community towards LLM chatbots?

Prompts serve as the primary interface between users and LLMs [39, 48]; however, little is known about how specific communities construct them.

We hypothesize that:

- **H1:** In everyday interactions with LLM chatbots, LGBTQ+ users rarely disclose their identity explicitly.
- **H2:** Prompting and LLM chatbot interactions from LGBTQ+ users are distinguishable from those of cis heterosexual users.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

Our analysis of distinguishability suggests... [ADD QUANT RESULT]

Our qualitative reading indicates... [ADD QUAL RESULT]

The key contributions from this work are:

- **C1:** An annotated dataset (COMMUNITY-SUBTEXT) of ecologically valid, naturally occurring prompts from members of the LGBTQ+ community and cisgender heterosexual baseline (Section ??).
- **C2:** An analysis of the features that distinguish LGBTQ+ prompts, including lexical, syntactic, and semantic characterizations (Section 4).
- **C3:** Measurements of LLM chatbot usage frequency, aversion, and satisfaction within the LGBTQ+ community (Section 4).

By centering the real prompting patterns of the LGBTQ+ community, our work informs the design of LLM chatbots that mitigate social bias across the full spectrum of identity expression.

2 Background

2.1 Prompting as interface

Prompting refers to the act of providing text as input to an LLM. Although commonly understood as giving instructions to guide model behavior [67], LLMs can ingest any possible finite string that fits within their context window, including confessions, musings, rants, and prayers. In many ways, LLMs 'react' to human expression rather than strictly 'follow orders'.

Human-Computer Interaction (HCI) has previously studied whether users can prompt effectively [24, 55, 68]. We take the field into a different set of questions: What do prompts reveal and conceal about the people who write them? People report that their TikTok algorithm knew they were part of the LGBTQ+ community even before they did. [

How to convey the significance of prompt data for HCI?]

[14, 50, 52, 67, 71]

Queer HCI [57]. Some queer people resist GenAI [56].

2.2 Implicit language signals

Implicit language signals affect how LLMs behave towards users [23, 25]. The effect of explicit signals is much better understood, especially because it is easier to implement in evaluations. Although explicit signals still account for a lot of undesired behavior, safety and behavioral control mechanisms are more effective around explicit signals than implicit signals.

[13] [3] [16] [27]

What complicates things is that implicit signals might not even map to easily identifiable characteristics [60]. The hidden variables could be closer to *vibes* than to discrete and clear identity classes.

2.3 Lavender linguistics

Lavender linguistics [30]

Past studies have tried to identify gay male speech, but nothing re: written text? prompts? **[Anything we want to say about this?]** [2, 9, 10, 22, 28, 32, 41, 66]

2.4 Queer legibility

There is a lot to say about queerness and **legibility**. There are several factors that complicate how LGBTQ+ people are read.

In the closet. Many LGBTQ+ people live in fear. Survival depends on the refusal to be decoded. Effort to hide any distinguishing language. The signal is squashed.

Cruising on the DL. Discreet communication. Just enough gesture. In-group recognition but otherwise extreme discretion. The signal is carefully modulated.

At the rave and the ball. The *surround* and the underground are sites of aesthetic, cultural, and intellectual production. Right to both opacity and identification. Some ephemera become LGBTQ+ tradition. Stereotypes don't reflect the party inside. The signal hits resonance.

At the Pride parade and on Netflix. Pinkwashing and the mainstream. Style extraction. The signal is amplified and diluted.

Towards utopia. Not a stable code. New slang and rhythm. New frequencies are in the air. New signals are echoing from the future.

[21] [28] [15]. [40] [64] [65]

2.5 Related Work

Although previous work has explored how vulnerable communities are treated differently by GenAI in both explicit and implicit settings, we identify the following gaps:

- Lack of focus on LGBTQ+ community [23, 26, 36, 61].
- Use of synthetic prompts [12, 63], repurposed texts [43, 51, 59, 62] or artificial injections [54, 58] instead of real user prompts.

What distinguishes our work from prior efforts is the intersection of **LGBTQ+ focus** and **ecological validity**: the degree to which the stimuli and experimental conditions reflect how the system is used in practice [8]. Our work is a response to the growing body of literature advocating for greater ecological validity in LLM evaluations [4, 19, 31].

3 Methods

3.1 Data collection

[Nolan?]

Refer to Appendix C for full survey design.

3.2 Data annotation

3.2.1 *Explicit.* We annotate **markedness** by...

[LLM Judges good for this?]

3.2.2 *Implicit.* We perform community-driven [64] annotation of **codedness** by...

[Community human annotation is best! Another survey to annotate or just us?]

3.3 Prompt dataset

Details of dataset are in Appendix A

3.4 Distinguishability analysis

We analyze how distinguishable the LGBTQ+ prompt set is from the baseline cis heterosexual prompt set through a variety of lenses.

3.4.1 Lexical. We quantify the extent to which vocabulary differs across prompt sets by applying *Calibrated Marked Words* from [38] with an added Benjamini–Hochberg (BH) procedure [5].

3.4.2 Syntactic. We measure how much the syntactic structure differs across prompt sets by generating syntactic labels for every prompt using NeuroBiber [1]. Unlike the open lexical vocabulary, NeuroBiber yields a small, fixed inventory of binary features, so calibration is unnecessary, and we use the simple log-odds ratio.

3.4.3 Semantic. We would like to determine if both prompt sets cover the same range of meaning. We embed each prompt into a shared semantic space and use a variant of the Maximum Mean Discrepancy (MMD-Fuse [7]) to assess whether the resulting embedding sets are drawn from different distributions. Our semantic embeddings are the Mean-pool residual stream at 75% of the depth. We consider Qwen3.6 27B [47], Llama 3.3 70B [37] and Gemma 4 31B [20].

3.4.4 Distributional. To estimate how separable the prompt sets are in principle, we perform a Classifier Two-Sample Test (C2ST) [34]. The held-out accuracy lower-bounds the Bayes-optimal separability. We consider a linear classifier on the semantic embeddings of Section 3.4.3.

3.4.5 Attributional. We identify the concept that most distinguishes LGBTQ+ prompts by attributing the probe’s decision to individual tokens. The distributional probe (Section 3.4.4) applies a linear head $\mathbf{w}$ to the mean-pooled residual stream, so its score decomposes exactly into a sum of per-token contributions $a_t = \frac{1}{T} \mathbf{w}^\top \mathbf{h}_t$, where $\mathbf{h}_t$ is the residual for token t and T the prompt length. A positive contribution moves the probe toward the LGBTQ+ class and a negative one toward the baseline. Because these contributions sum to the probe score, they are exact rather than estimated. We rank prompts by the tokens that contribute most to obtaining the largest number of LGBTQ+-marked texts, along with a contrasting set of non-activating texts. Claude Opus 4.8 reads both sets, with the strongest-contributing tokens highlighted inline, and names the discriminating concept.

3.4.6 Topical. **[Should we define orthogonal topics?]** We assess whether LGBTQ+ users prompt about different subjects than the cis-heterosexual baseline. We assign a topic to each prompt using TopicGPT [45], yielding topic distributions over the LGBTQ+ and baseline sets, and quantify the difference between them using the Jensen–Shannon divergence.

[Or instead, measure real vs. hypothetical in already-defined survey topics?]

3.4.7 Interactional. **[Brainstorm:]**

[Speech act annotation: information-seeking vs. advice-seeking vs. emotional disclosure vs. venting]
 [Self-disclosure depth: none vs. factual vs. emotional disclosure.]
 [anthropomorphization]

3.5 Qualitative analysis of single prompts

We open-code the prompts in the grounded-theory tradition, letting codes emerge from the text rather than imposing them. Two annotators independently code a subsample, reconcile a shared codebook, and code the remainder to saturation.

3.6 Multi-turn analysis of donated transcripts

Donated transcripts give us multi-turn dynamics.

Forking turns. Identify the point in the transcript at which the conversation takes a sharp turn and forks off [6, 29, 72].

This could include:

- Prompt where markedness or codedness is introduced by user
- Prompt after which the LLM showed blatant bias, shift in behavior, or user reported feeling harmed

4 Results

[just placeholder/ideations for now]

4.1 LGBTQ+ prompts do not often explicitly disclose identity

4.2 Codedness is very variable

4.3 Implicit signal present but no ideal *gaydar*

4.4 Vocabulary is the biggest differentiator

4.5 Trans and BIPOC resist AI more

4.6 Cisgay men have good experiences with AI

5 Discussion

[just placeholder/ideations for now]

Our central findings suggest that ...

Involuntary legibility. ...

Conditional distinguishability. Unknown hidden variables...

5.1 Limitations

Although our dataset is an advancement towards ecologically valid evaluations for LLM bias against LGBTQ+, there are some limitations:

- Mostly single-turn. [11]
- Not a wide international or multilingual representation
- Static snapshot
- Annotation subjectivity and scale

5.2 Future work

LLM behavioral study. Do distinguishable prompts elicit different behavior from LLMs?

5.3 Wider applications

Being able to distinguish prompt sets is a generally useful capability in prompt engineering as it allows for prompt filtering and quality prediction [70]

In particular, *prompt routing* could serve as a possible implementation of Pluralistic Alignment [53] by selecting personalized models per community [18].

The LGBTQ+ community is diverse and multiplicitous, constantly evolving and evading capture. LGBTQ+ prompts thus push the limits of decodability, testing how well LLMs can single out users by community.

6 Conclusion

This paper closes the gap between how bias is typically measured and how real-life LGBTQ+ identity is expressed.

Impact and ethics statement

Our work handles sensitive data from the LGBTQ+ community. Appropriate measures were taken to anonymize data before release.

In the wrong hands, distinguishability technology could be used against LGBTQ+ people, for surveillance and targeted discrimination.

Positionality statement

Our research team consists of members of the LGBTQ+ community from three different continents.

LLM usage statement

Judge LLMs were used to develop *markedness* annotations (Described in Section 3.2). TopicGPT was used to develop topic annotations (Described in Section 3.4.6). Claude Opus 4.8 was used for the conceptual attribution 3.4.5). Claude code was used as an initial aid in prototyping experiments. All output, code, and text have been verified by the authors.

References

- [1] Kenan Alkiek, Anna Wegmann, Jian Zhu, and David Jurgens. 2025. Neurobiber: Fast and Interpretable Stylistic Feature Extraction. arXiv:2502.18590 [cs.CL] <https://arxiv.org/abs/2502.18590>
- [2] Paul Baker. 2013. From gay language to normative discourse: A diachronic corpus analysis of Lavender Linguistics conference abstracts 1994–2012. *Journal of Language and Sexuality* 2, 2 (2013), 179–205. doi:10.1075/jls.2.2.01bak
- [3] Jeffrey Basoah, Daniel Chechelnitsky, Tao Long, Katharina Reinecke, Chrysoula Zerva, Kaitlyn Zhou, Mark Diaz, and Maarten Sap. 2025. Not Like Us, Hunt: Measuring Perceptions and Behavioral Effects of Minoritized Anthropomorphic Cues in LLMs (*FAccT '25*). Association for Computing Machinery, New York, NY, USA, 710–745. doi:10.1145/3715275.3732045
- [4] Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, Jan Batzner, Negar Foroutan, Chris Schmitz, Karolina Korgul, Hunar Batra, Oishi Deb, Emma Beharry, Cornelius Emde, Thomas Foster, Anna Gausen, Maria Grandury, Simeng Han, Valentin Hofmann, Lujain Ibrahim, Hazel Kim, Hannah Rose Kirk, Fangru Lin, Gabrielle Kaili-May Liu, Lennart Luetzgau, Jabez Magomere, Jonathan Rystrom, Anna Sotnikova, Yushi Yang, Yilun Zhao, Adel Bibi, Antoine Bosselut, Ronald Clark, Arman Cohan, Jakob Foerster, Yarin Gal, Scott A. Hale, Inioluwa Deborah Raji, Christopher Summerfield, Philip H. S. Torr, Cozmin Ududec, Luc Rocher, and Adam Mahdi. 2025. Measuring what Matters: Construct Validity in Large Language Model Benchmarks. arXiv:2511.04703 [cs.CL] <https://arxiv.org/abs/2511.04703>
- [5] Yoav Benjamini and Yosef Hochberg. 1995. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal statistical society: series B (Methodological)* 57, 1 (1995), 289–300.
- [6] Eric Bigelow, Ari Holtzman, Hidenori Tanaka, and Tomer Ullman. 2024. Forking Paths in Neural Text Generation. arXiv:2412.07961 [cs.CL] <https://arxiv.org/abs/2412.07961>

- [7] Felix Biggs, Antonin Schrab, and Arthur Gretton. 2023. MMD-FUSE: Learning and Combining Kernels for Two-Sample Testing Without Data Splitting. arXiv:2306.08777 [stat.ML] <https://arxiv.org/abs/2306.08777>
- [8] Egon Brunswik. 1956. *Perception and the Representative Design of Psychological Experiments* (2nd ed.). University of California Press, Berkeley, CA.
- [9] Jeremy Calder. 2020. Language and sexuality: Language and LGBTQ+ communities. *The international encyclopedia of linguistic anthropology* (2020), 1–7.
- [10] Jovelyn M Cantina. 2020. Gay argot: Beyond the coded words and meanings of lavender lexicon. *International Review of Humanities and Scientific Research* (2020), 248–262.
- [11] Mohit Chandra, Siddharth Sriraman, Harneet Singh Khanuja, Yiqiao Jin, and Munmun De Choudhury. 2026. Reasoning Is Not All You Need: Examining LLMs for Multi-Turn Mental Health Conversations. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens (Eds.). Association for Computational Linguistics, San Diego, California, United States, 46648–46682. doi:10.18653/v1/2026.acl-long.2164
- [12] Crystal T Chang, Neha Srivathsa, Charbel Bou-Khalil, Akshay Swaminathan, Mitchell R Lunn, Kavita Mishra, Sanmi Koyejo, and Roxana Daneshjou. 2025. Evaluating anti-LGBTQIA+ medical bias in large language models. *PLOS Digital Health* 4, 9 (2025), e0001001.
- [13] S. Chapagain, Y. Zhao, T. K. Rohleen, S. M. Hamdi, S. F. Boubrahimi, R. E. Flinn, E. M. Lund, D. Klooster, J. R. Scheer, and C. J. Cascalheira. 2024. Predictive Insights into LGBTQ+ Minority Stress: A Transductive Exploration of Social Media Discourse. arXiv:2411.13534 [cs.CL] <https://arxiv.org/abs/2411.13534>
- [14] Hai Dang, Sven Goller, Florian Lehmann, and Daniel Buschek. 2023. Choice Over Control: How Users Write with Large Language Models using Diegetic and Non-Diegetic Prompting. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. ACM, 1–17. doi:10.1145/3544548.3580969
- [15] Nicholas de Villiers. 2012. *Opacity and the Closet: Queer Tactics in Foucault, Barthes, and Warhol*. University of Minnesota Press, Minneapolis.
- [16] Rebecca Dorn, Lee Kezar, Fred Morstatter, and Kristina Lerman. 2024. Harmful Speech Detection by Language Models Exhibits Gender-Queer Dialect Bias. arXiv:2406.00020 [cs.CL] <https://arxiv.org/abs/2406.00020>
- [17] Dget L. Downey and David A. Jenkins. 2026. AI in Supporting LGBTQIA+ Populations. In *Artificial Intelligence in Social Work*, Ruopeng An and Michael A. Lindsey (Eds.). Springer, Cham, 125–147. doi:10.1007/978-3-032-18443-6_7
- [18] Shangbin Feng, Taylor Sorensen, Yuhan Liu, Jillian Fisher, Chan Young Park, Yejin Choi, and Yulia Tsvetkov. 2024. Modular Pluralism: Pluralistic Alignment via Multi-LLM Collaboration. arXiv:2406.15951 [cs.CL] <https://arxiv.org/abs/2406.15951>
- [19] Bufan Gao and Elisa Kreiss. 2025. Measuring Bias or Measuring the Task: Understanding the Brittle Nature of LLM Gender Biases. *arXiv preprint arXiv:2509.04373* (2025). arXiv:2509.04373 [cs.CL] <https://arxiv.org/abs/2509.04373>
- [20] Gemma Team, Google DeepMind. 2026. Gemma 4: Byte for Byte, the Most Capable Open Models. <https://blog.google/innovation-and-ai/technology/developers-tools/gemma-4/>. Model documentation: <https://ai.google.dev/gemma/docs/core>.
- [21] Édouard Glissant. 1997. *Poetics of Relation*. University of Michigan Press, Ann Arbor. Trans. Betsy Wing; orig. *Poétique de la Relation*, Gallimard, 1990. See “For Opacity”.
- [22] Kaikoura Gutteridge. 2024. The Gay Way to Say: A Review of Sociolinguistic Research on Gay Speech. *Journal of Undergraduate Linguistics* 3, 1 (2024), 92–100.
- [23] Irti Haq and Belén Saldías. 2026. Dialect vs Demographics: Quantifying LLM Bias from Implicit Linguistic Signals vs. Explicit User Profiles. arXiv:2604.21152 [cs.CL] <https://arxiv.org/abs/2604.21152>
- [24] Zeyu He, Sania Naphade, and Ting-Hao ‘Kenneth’ Huang. 2025. Prompting in the Dark: Assessing Human Performance in Prompt Engineering for Data Labeling When Gold Labels Are Absent. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. ACM, New York, NY, USA. doi:10.1145/3706598.3714319
- [25] Valentin Hofmann, Pratyusha Ria Kalluri, Dan Jurafsky, and Sharese King. 2024. AI generates covertly racist decisions about people based on their dialect. *Nature* 633, 8028 (2024), 147–154.
- [26] Kyuha Jung, Gyuho Lee, Yuanhui Huang, and Yunan Chen. 2025. ‘I’ve talked to ChatGPT about my issues last night’: Examining Mental Health Conversations with Large Language Models through Reddit Analysis. *Proceedings of the ACM on Human-Computer Interaction* 9, 7 (2025), 1–25.
- [27] Amir Karami, Morgan Lundy, Frank Webb, Hannah R Boyajieff, Michael Zhu, and Dorathea Lee. 2021. Automatic categorization of LGBT user profiles on Twitter with machine learning. *Electronics* 10, 15 (2021), 1822.
- [28] Don Kulick. 2000. Gay and Lesbian Language. *Annual Review of Anthropology* 29 (2000), 243–285. doi:10.1146/annurev.anthro.29.1.243
- [29] Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. 2025. LLMs Get Lost In Multi-Turn Conversation. arXiv:2505.06120 [cs.CL] <https://arxiv.org/abs/2505.06120>
- [30] William L Leap. 2015. Lavender language. *The International Encyclopedia of Human Sexuality* (2015), 649–719.
- [31] Charlotte Li, Nick Hagar, Sachita Nishal, Jeremy Gilbert, and Nick Diakopoulos. 2026. Towards Real-World Validity in Generative AI Benchmarks: Understanding and Designing Domain-Centered Evaluations for Journalism Practitioners. arXiv:2511.05501 [cs.HC] <https://arxiv.org/abs/2511.05501>
- [32] Aaron Loh, Kenneth Soo, and Hui-lin Xing. 2016. Predicting Sexual Orientation based on Facebook Status. *Projects in the Course CS229: Machine Learning* (2016).
- [33] Yinru Long, Zilin Ma, Yiyang Mei, and Zhaoyuan Su. 2024. AffirmativeAI: Towards LGBTQ+ Friendly Audit Frameworks for Large Language Models. In *1st HEAL Workshop at CHI Conference on Human Factors in Computing Systems*. Honolulu, HI, USA. arXiv:2405.04652 [cs.HC] <https://arxiv.org/abs/2405.04652>

- [34] David Lopez-Paz and Maxime Oquab. 2018. Revisiting Classifier Two-Sample Tests. arXiv:1610.06545 [stat.ML] <https://arxiv.org/abs/1610.06545>
- [35] Zilin Ma, Yiyang Mei, Yinru Long, Zhaoyuan Su, and Krzysztof Z. Gajos. 2024. Evaluating the Experience of LGBTQ+ People Using Large Language Model-Based Chatbots for Mental Health Support. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)* (Honolulu, HI, USA). Association for Computing Machinery, New York, NY, USA, 1–15. arXiv:2402.09260 [cs.HC] doi:10.1145/3613904.3642482
- [36] Alba María Mármol-Romero, Adrián Moreno-Muñoz, Flor Miriam Plaza-del Arco, M. Dolores Molina-González, and Arturo Montejó-Ráez. 2024. MentalRiskES: A New Corpus for Early Detection of Mental Disorders in Spanish. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*. ELRA and ICCL, Torino, Italia, 11204–11214. <https://aclanthology.org/2024.lrec-main.978>
- [37] Meta AI. 2024. Llama 3.3 70B Model Card. https://github.com/meta-llama/llama-models/blob/main/models/llama3_3/MODEL_CARD.md. Accessed: 2026-07-06.
- [38] Jennifer Mickel, Maria De-Arteaga, Leqi Liu, and Kevin Tian. 2025. More of the Same: Persistent Representational Harms Under Increased Representation. arXiv:2503.00333 [cs.CL] <https://arxiv.org/abs/2503.00333>
- [39] Meredith Ringel Morris. 2024. Prompting Considered Harmful. *Commun. ACM* (2024).
- [40] José Esteban Muñoz. 2009. *Cruising Utopia: The Then and There of Queer Futurity*. New York University Press, New York.
- [41] Benjamin Munson, Elizabeth C. McDonald, Nancy L. DeBoe, and Aubrey R. White. 2006. The acoustic and perceptual bases of judgments of women and men’s sexual orientation from read speech. *Journal of Phonetics* 34, 2 (2006), 202–240. doi:10.1016/j.wocn.2005.05.003
- [42] Ruby Ostrow and Adam Lopez. 2025. LLMs Reproduce Stereotypes of Sexual and Gender Minorities. arXiv:2501.05926 [cs.CL] <https://arxiv.org/abs/2501.05926>
- [43] Anaelia Ovalle, Palash Goyal, Jwala Dhamala, Zachary Jaggars, Kai-Wei Chang, Aram Galstyan, Richard Zemel, and Rahul Gupta. 2023. “I’m fully who I am”: Towards Centering Transgender and Non-Binary Voices to Measure Biases in Open Language Generation. In *2023 ACM Conference on Fairness, Accountability and Transparency (FAccT ’23)*. ACM, 1246–1266. doi:10.1145/3593013.3594078
- [44] Srikant Panda, Hitesh Laxmichand Patel, Shahad Al-Khalifa, Amit Agarwal, Hend Al-Khalifa, and Sharefah Al-Ghamdi. 2025. DAIQ: Auditing Demographic Attribute Inference from Question in LLMs. *arXiv preprint arXiv:2508.15830* (2025).
- [45] Chau Minh Pham, Alexander Hoyle, Simeng Sun, Philip Resnik, and Mohit Iyyer. 2024. TopicGPT: A Prompt-based Topic Modeling Framework. arXiv:2311.01449 [cs.CL] <https://arxiv.org/abs/2311.01449>
- [46] Elinor Poole-Dayana, Deb Roy, and Jad Kabbara. 2026. LLM targeted underperformance disproportionately impacts vulnerable users. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 40. 39116–39124.
- [47] Qwen Team. 2026. Qwen3.6. <https://github.com/QwenLM/Qwen3.6>. Open-weight release (Qwen3.6-35B-A3B, Qwen3.6-27B). Accessed 2026-07-06.
- [48] Laria Reynolds and Kyle McDonell. 2021. Prompt Programming for Large Language Models: Beyond the Few-Shot Paradigm. In *Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems (CHI EA ’21)*. ACM, New York, NY, USA. doi:10.1145/3411763.3451760
- [49] Ian Rios-Sialer. 2026. The Homogenization Problem in LLMs: Towards Meaningful Diversity in AI Safety. arXiv:2601.06116 [cs.AI] <https://arxiv.org/abs/2601.06116>
- [50] Rupak Sarkar, Bahareh Sarrafzadeh, Nirupama Chandrasekaran, Nagu Rangan, Philip Resnik, Longqi Yang, and Sujay Kumar Jauhar. 2025. Conversational User-AI Intervention: A Study on Prompt Rewriting for Improved LLM Response Generation. arXiv:2503.16789 [cs.CL] <https://arxiv.org/abs/2503.16789>
- [51] Morgan Klaus Scheuerman, Katy Weathington, Adrian Petterson, Dylan Thomas Doyle, Dipto Das, Michael Ann DeVito, and Jed R. Brubaker. 2025. Transphobia Is in the Eye of the Prompter: Trans-Centered Perspectives on Large Language Models. *ACM Transactions on Computer-Human Interaction* 32, 5 (2025), 1–42. doi:10.1145/3743676
- [52] Yuanfeng Song, Yuanqin He, Xuefang Zhao, Hanlin Gu, Di Jiang, Haijun Yang, Lixin Fan, and Qiang Yang. 2023. A Communication Theory Perspective on Prompting Engineering Methods for Large Language Models. arXiv:2310.18358 [cs.CL] <https://arxiv.org/abs/2310.18358>
- [53] Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, Tim Althoff, and Yejin Choi. 2024. A Roadmap to Pluralistic Alignment. arXiv:2402.05070 [cs.AI] <https://arxiv.org/abs/2402.05070>
- [54] Mae Sosto, Delfina Sol Martinez Pandiani, and Laura Hollink. 2026. QueerGen: How LLMs Reflect Societal Norms on Gender and Sexuality in Sentence Completion Tasks. *arXiv preprint arXiv:2601.20731* (2026). arXiv:2601.20731 [cs.CL] <https://arxiv.org/abs/2601.20731>
- [55] Hari Subramonyam, Christopher Lawrence Pondoc, Colleen Seifert, Maneesh Agrawala, and Roy Pea. 2024. Bridging the Gulf of Envisioning: Cognitive Challenges in Prompt Based Interactions with LLMs. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems (CHI ’24)*. ACM, New York, NY, USA. arXiv:2309.14459.
- [56] Jordan Taylor, Joel Mire, Alicia DeVrio, Maarten Sap, Haiyi Zhu, and Sarah E. Fox. 2026. “I Just Don’t Want My Work Being Fed Into The AI Blender”: Queer Artists on Refusing and Resisting Generative AI. arXiv:2604.14266 [cs.HC] <https://arxiv.org/abs/2604.14266>
- [57] Jordan Taylor, Ellen Simpson, Anh-Ton Tran, Jed R. Brubaker, Sarah E Fox, and Haiyi Zhu. 2024. Cruising Queer HCI on the DL: A Literature Review of LGBTQ+ People in HCI. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI ’24)*. ACM, 1–21. doi:10.1145/3613904.3642494
- [58] Alberto Testoni and Iacer Calixto. 2026. Calibrated? Not for Everyone: How Sexual Orientation and Religious Markers Distort LLM Accuracy and Confidence in Medical QA. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens (Eds.). Association for Computational Linguistics, San Diego, California, United States, 432–447. doi:10.18653/v1/2026.acl-short.36

- [59] Joshua Tint. 2025. Guardrails, not Guidance: Understanding Responses to LGBTQ+ Language in Large Language Models. In *Proceedings of the Queer in AI Workshop*, A Pranav, Alissa Valentine, Shaily Bhatt, Yanan Long, Arjun Subramonian, Amanda Bertsch, Anne Lauscher, and Ankush Gupta (Eds.). Association for Computational Linguistics, Hybrid format (in-person and virtual), 6–16. doi:10.18653/v1/2025.queerina-main.2
- [60] Nenad Tomasev, Kevin R McKee, Jackie Kay, and Shakir Mohamed. 2021. Fairness for unobserved characteristics: Insights from technological impacts on queer communities. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*. 254–265.
- [61] Manuel Tonneau, Neil K. R. Seghal, Niyati Malhotra, Sharif Kazemi, Victor Orozco-Olvera, Ana Maria Muñoz Boudet, Lakshmi Subramanian, Samuel P. Fraiberger, Sharath Chandra Guntuku, and Valentin Hofmann. 2026. Different Demographic Cues Yield Inconsistent Conclusions About LLM Personalization and Bias. arXiv:2601.18486 [cs.CL] <https://arxiv.org/abs/2601.18486>
- [62] Juan Vásquez, Gemma Bel-Enguix, Scott Thomas Andersen, and Sergio-Luis Ojeda-Trueba. 2022. HeteroCorpus: A Corpus for Heteronormative Language Detection. In *Proceedings of the 4th Workshop on Gender Bias in Natural Language Processing (GeBNLP)*, Christian Hardmeier, Christine Basta, Marta R. Costa-jussà, Gabriel Stanovsky, and Hila Gonen (Eds.). Association for Computational Linguistics, Seattle, Washington, 225–234. doi:10.18653/v1/2022.gebnlp-1.23
- [63] Bhaskara Hanuma Vedula, Darshan Anghan, Ishita Goyal, Ponnuram Kumaraguru, and Abhijnan Chakraborty. 2026. ImplicitBBQ: Benchmarking Implicit Bias in Large Language Models through Characteristic-Based Cues. arXiv:2604.01925 [cs.CL] <https://arxiv.org/abs/2604.01925>
- [64] Leonor Veloso, Lea Hirilmann, Philipp Wicke, and Hinrich Schütze. 2025. SLAYiNG: Towards Queer Language Processing. arXiv:2509.17449 [cs.CL] <https://arxiv.org/abs/2509.17449>
- [65] McKenzie Wark. 2023. *Raving*. Duke University Press.
- [66] Hsiao-Han Wu and Shu-Kai Hsieh. 2017. Exploring Lavender Tongue from Social Media Texts [In Chinese]. In *Proceedings of the 29th Conference on Computational Linguistics and Speech Processing (ROCLING 2017)*. 68–80.
- [67] Liangwei Yang, Shiyu Wang, Haolin Chen, Rithesh Murthy, Ming Zhu, Jielin Qiu, Zixiang Chen, Juntao Tan, Jianguo Zhang, Zhiwei Liu, Wenting Zhao, Silvio Savarese, Caiming Xiong, Huan Wang, and Shelby Heinecke. 2026. Position: Vector Prompt Interfaces Should Be Exposed to Enable Customization of Large Language Models. arXiv:2603.04292 [cs.CL] <https://arxiv.org/abs/2603.04292>
- [68] J.D. Zamfirescu-Pereira, Richmond Y. Wong, Bjoern Hartmann, and Qian Yang. 2023. Why Johnny Can’t Prompt: How Non-AI Experts Try (and Fail) to Design LLM Prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI ’23)*. ACM, New York, NY, USA. doi:10.1145/3544548.3581388
- [69] Marc Zao-Sanders. 2026. How People Are Really Using AI in 2026. *Harvard Business Review* (1 June 2026). <https://hbr.org/2026/06/how-people-are-really-using-ai-in-2026> Digital article, Generative AI section.
- [70] Yuanming Zhang, Yan Lin, Arijit Khan, and Huaiyu Wan. 2026. Large Language Model Prompt Datasets: An In-depth Analysis and Insights. arXiv:2510.09316 [cs.LG] <https://arxiv.org/abs/2510.09316>
- [71] Shengqi Zhu, Jeffrey M. Rzeszotarski, and David Mimno. 2026. Priming, Path-dependence, and Plasticity: Understanding the molding of user-LLM interaction and its implications from (many) chat logs in the wild. arXiv:2605.05767 [cs.HC] <https://arxiv.org/abs/2605.05767>
- [72] Amir Zur, Atticus Geiger, Ekdeep Singh Lubana, and Eric Bigelow. 2025. Are language models aware of the road not taken? Token-level uncertainty and hidden state dynamics. arXiv:2511.04527 [cs.CL] <https://arxiv.org/abs/2511.04527>

Appendix A Datasets

A.1 COMMUNITY-SUBTEXT overview

We built COMMUNITY-SUBTEXT: a dataset of prompts that users pose to LLMs about mental health, gender/sexual health, and relationships. Each prompt is drawn from the survey (Section 3.1), and annotated (Section 3.2) for how that identity surfaces in the text:

$$\mathcal{D} = \{ (x_i, y_i, z_i, d_i, c_i) \}_{i=1}^n.$$

Here x_i is the prompt and each attribute group serves a distinct role:

- y — **labels**: Coarse group membership and annotations of the prompt;
- z — **identity**: Detailed self-reported identity of the author;
- d — **demographics** of the author;
- c — **interaction context** from the survey.

Throughout, * marks a prefer-not-to-answer response.

Labels, identity, and demographics (Table A.1). The coarse labels y are the main signal to differentiate prompts of the LGBTQ+ community from a cisheterosexual baseline; z records the author’s detailed self-reported LGBTQ+ identity. The survey also recorded a wide range of self-reported demographic information d .

Interaction context (Table A.2). We generally characterize attitude and frequency through the ordinal scales $\mathcal{A}$ and $\mathcal{F}$. Each topic maps to one domain, with five topics per domain: MH = mental health, GSH = gender/sexual health, REL = relationships.

Table A.1. Per-prompt annotations: coarse group labels y , detailed LGBTQ+ identity z , and demographics d of the author. Multi-select fields are subsets ($\subseteq$) of their option sets; * marks a prefer-not-to-answer response.

(a) Coarse group labels				(b) Detailed LGBTQ+ identity		
Field	Symbol	Values	Meaning	Field	Symbol	Values
LGBTQ+ label	y^{lgbtq}	$\{0, 1\}$	LGBTQ+ vs. baseline	Transgender	z^{trans}	$\{0, 1, *\}$
Markedness	y^{mark}	$\{0, 1\}$	explicit signal	Gender	z^{gen}	$\subseteq \{\text{woman, man, nonbinary, other}\}$
Codedness	y^{code}	$[0, 1]$	implicit signal	Orientation	z^{ori}	$\subseteq \{\text{asexual, bisexual, gay, lesbian, pansexual, queer, questioning, straight, other}\}$
				Pronouns	z^{pro}	$\subseteq \{\text{she/her, he/him, they/them, other}\}$
(c) Demographics						
Field	Symbol	Values				
Race	d^{race}	multi-select $\subseteq \{\text{Hispanic, Latino, Black/African/African American, Central Asia, North Asia, South Asia, Southeast Asia, West Asia, Native American, Middle Eastern, Pacific Islander, White/Caucasian, other}\}$				
Age	d^{age}	ordinal: 18–24 < 25–34 < 35–44 < 45–54 < 55–64 < 65–74 < 75–84 < 85–94 < 95+				
Disability	d^{dis}	$\{0, 1, *\}$				
Education	d^{edu}	ordinal: <high school < high school < some college < associate < bachelor < master < professional < doctorate				
Income	d^{inc}	ordinal: none < \$1–10k < \$10–25k < \$25–50k < \$50–100k < \$100–125k < \$125–150k < \$150–175k < \$175–200k < \$200k+				

Table A.2. Interaction context c . Left: fields recorded per prompt. Right: the attitude scale $\mathcal{A}$ and frequency scale $\mathcal{F}$. Bottom: the fifteen topics, grouped by domain.

(a) Fields			(b) Ordinal scales		
Field	Symbol	Values	Level	$\mathcal{A}$ (attitude)	$\mathcal{F}$ (frequency)
Topic	c^{top}	$\{1, \dots, 15\}$	1	very negative	very rarely
Domain	c^{dom}	$\{\text{MH, GSH, REL}\}$	2	somewhat negative	~once/yr
Provenance	c^{prov}	$\{\text{real, hyp}\}$	3	neutral	~once/mo
Adoption	c^{hist}	ordinal, 5 lvl	4	somewhat positive	~once/wk
General freq.	c^{gen}	$\mathcal{F}$	5	very positive	~once/day
LLM freq.	c^{freq}	$\mathcal{F}$	6	—	1–5/day
Professional freq.	c^{prof}	$\mathcal{F}$	7	—	5–15/day
Aversion	c^{avr}	$\mathcal{A}$	8	—	>15/day
Satisfaction	c^{sat}	$\mathcal{A}$			
(c) Topics (c^{top}) by domain (c^{dom})					
MH — mental health		GSH — gender/sexual health		REL — relationships	
1	isolation, anxiety, depression, or panic in social settings	6	understanding one's gender or sexual orientation	11	communication, intimacy, or boundaries in a new relationship
2	coping with a major life transition or shift in self-view	7	STI symptoms, causes, or treatments	12	exploring a new relationship dynamic
3	whether to start seeing a therapist	8	understanding gender-affirming care	13	communicating needs to a partner
4	values conflicting with family or community expectations	9	side effects of gender/sexual-health medications (PrEP, estrogen)	14	media featuring relationships like one's own
5	improving body image	10	finding knowledgeable healthcare providers	15	preparing for a new sexual experience

A.2 SUBJECT-TRAJECTORIES overview

The dataset is available ... We also built SUBJECT-TRAJECTORIES: a multi-turn dataset derived from donated transcripts and qualitative labels per turn:

$$\mathcal{T} = \left\{ (x_0, y_0, x_1, y_1, \dots)_i \right\}_{i=1}^n.$$

A.3 Dataset access

The dataset is available ... **[is anonymized dataset good to share publicly? or do MentalRiskES-like type gating?]**

Appendix B Distinguishing other groups

To validate our pipeline (Section 3.4), we apply it to publicly available datasets to validate that our methods capture distinguishability in other domains.

Appendix C Survey Design

This appendix reproduces the complete survey instrument (*Survey on the Use of LLMs for Mental Health, Gender, and Sexuality Advice*, STUDY2026_00000255). Items marked * were required; response types are annotated in gray. Because Sections C.5–C.7 share one branching template that varies only in the *question topic* (blue), the template is given in full once (Appendix C.4); per-item topics are then listed as substitutions.

C.1 Survey Description (shown to participants)

Participants draw on personal experience to describe how they naturally interact with LLMs / AI chatbots for support regarding mental health, gender, and sexuality. The stated goal is a community-centered dataset of authentic interactions to measure bias in AI systems. Participants may optionally upload complete chatbot transcript data for a \$5 bonus, and are asked to answer authentically without revealing private or personally identifiable information.

C.2 Pre-Survey Questions

[History] When did you start regularly using AI chatbots?

- Within the past year
- Within the past 1–2 years
- Within the past 2–5 years
- More than 5 years ago
- I have never regularly used AI chatbots

[General Frequency] In general, how often do you use an AI chatbot?

- Very rarely
- About once a year
- About once a month
- About once a week
- About once a day
- About 1–5 times a day
- About 5–15 times a day

- More than 15 times a day

[Frequency: LLM Mental Health] How often do you use an AI chatbot for advice about mental health?

- Very rarely
- About once a year
- About once a month
- About once a week
- About once a day
- About 1–5 times a day
- About 5–15 times a day
- More than 15 times a day

[Frequency: Other Sources Mental Health] How often do you seek professional advice about mental health (e.g., mental health professional, therapist, doctor, professional medical website)?

- Very rarely
- About once a year
- About once a month
- About once a week
- About once a day
- About 1–5 times a day
- About 5–15 times a day
- More than 15 times a day

[Frequency: LLM Gender/Sexual Health] How often do you use an AI chatbot for advice about gender or sexual health?

- Very rarely
- About once a year
- About once a month
- About once a week
- About once a day
- About 1–5 times a day
- About 5–15 times a day
- More than 15 times a day

[Frequency: Other Sources Gender/Sexual Health] How often do you seek professional advice about gender or sexual health (e.g., sexual health professional, endocrinologist, doctor, professional medical website)?

- Very rarely
- About once a year
- About once a month
- About once a week
- About once a day
- About 1–5 times a day

- 677 ○ About 5–15 times a day
- 678 ○ More than 15 times a day
- 679

680 **[Frequency: LLM Relationship Advice]** How often do you use an AI chatbot for advice about relationships?

- 681 ○ Very rarely
- 682 ○ About once a year
- 683 ○ About once a month
- 684 ○ About once a week
- 685 ○ About once a day
- 686 ○ About 1–5 times a day
- 687 ○ About 5–15 times a day
- 688 ○ More than 15 times a day
- 689
- 690
- 691

692 **[Frequency: Other Sources Relationship Advice]** How often do you seek professional advice about relationships
(e.g., therapist, counselor, professional counseling website)?

- 695 ○ Very rarely
- 696 ○ About once a year
- 697 ○ About once a month
- 698 ○ About once a week
- 699 ○ About once a day
- 700 ○ About 1–5 times a day
- 701 ○ About 5–15 times a day
- 702 ○ More than 15 times a day
- 703
- 704
- 705

706 **[Aversion: Mental Health]** How do you feel about the use of AI chatbots for advice about mental health?

- 707 ○ Very positive
- 708 ○ Somewhat positive
- 709 ○ Neutral
- 710 ○ Somewhat negative
- 711 ○ Very negative
- 712
- 713

714 **[Aversion: Gender & Sexual Health]** How do you feel about the use of AI chatbots for advice about gender or sexual
health?

- 717 ○ Very positive
- 718 ○ Somewhat positive
- 719 ○ Neutral
- 720 ○ Somewhat negative
- 721 ○ Very negative
- 722
- 723

724 **[Aversion: Relationship Advice]** How do you feel about the use of AI chatbots for advice about relationships?

- 725 ○ Very positive
- 726 ○ Somewhat positive
- 727
- 728

- Neutral
- Somewhat negative
- Very negative

C.3 Demographic Information

C.3.1 *Detailed Queer Identity Questions.* [Q] Do you identify as a member of the LGBTQ+ community?

- Yes
- No
- Prefer not to say

[Q] Do you identify as transgender?

- Yes
- No
- Prefer not to say

[Q] Which of the following best describes your gender? (*Select all that apply*)

- Woman
- Man
- Non-Binary
- Prefer not to say
- Prefer to self-describe [free-text response]

[Q] Do you identify as straight or heterosexual?

- Yes
- No
- Prefer not to say

[Q] Which of the following terms best describe your sexual orientation? (*Select all that apply*)

- Asexual
- Bisexual
- Gay
- Lesbian
- Pansexual
- Queer
- Questioning
- Straight or heterosexual
- Prefer to self-describe [free-text response]

[Q] What pronouns do you use? (*Used to address you in any future research publications; select all that apply.*)

- She/her/hers
- He/him/his

- 781 ○ They/them/theirs
- 782 ○ Prefer to self-describe [free-text response]
- 783

784
785 **C.3.2 Standard Demographic Questions. [Q] What is your age?**

- 786 ○ 18–24
- 787 ○ 25–34
- 788 ○ 35–44
- 789 ○ 45–54
- 790 ○ 55–64
- 791 ○ 65–74
- 792 ○ 75–84
- 793 ○ 85–94
- 794 ○ 95+
- 795
- 796
- 797

798 **[Q] Which racial or ethnic groups do you identify with? (*Select all that apply*)**

- 799 ○ Hispanic
- 800 ○ Latino
- 801 ○ Black / African / African American
- 802 ○ Central Asia
- 803 ○ North Asia
- 804 ○ South Asia
- 805 ○ Southeast Asia
- 806 ○ West Asia
- 807 ○ Native American
- 808 ○ Middle Eastern
- 809 ○ Pacific Islander
- 810 ○ White / Caucasian
- 811 ○ Prefer to self-describe [free-text response]
- 812 ○ I prefer not to answer this question
- 813
- 814
- 815
- 816
- 817

818 **[Q] Do you identify as a person with a disability?**

- 819 ○ Yes
- 820 ○ No
- 821 ○ I prefer not to answer this question
- 822
- 823

824 **[Q] What is your highest degree or level of education?**

- 825 ○ Less than a high school degree
- 826 ○ High school diploma or the equivalent
- 827 ○ Some college credit, no degree
- 828 ○ Associate degree
- 829 ○ Bachelor's degree
- 830
- 831
- 832

- Master's degree
- Professional degree beyond a bachelor's degree
- Doctorate degree
- I prefer not to answer this question

[Q] Approximate annual income in the past 12 months (*US Dollars*)?

- No income
- \$1 to \$10,000
- \$10,000 to \$24,999
- \$25,000 to \$49,999
- \$50,000 to \$99,999
- \$100,000 to \$124,999
- \$125,000 to \$149,999
- \$150,000 to \$174,999
- \$175,000 to \$199,999
- \$200,000 or more
- I prefer not to answer this question

C.4 Branching Question Template (Q1–Q15)

Every item in Sections C.5–C.7 follows the template below; only the *topic* changes. The template is shown with the *topic* of Q1.

[Qn] Have you ever turned to an AI chatbot *for advice on dealing with feelings of isolation, anxiety, depression, panic, or other mental health issues at work, school, or another social setting?**

↪ if Yes:

[Qn.1] What was the initial question/prompt you asked the AI chatbot to answer?* [free-text response]

[Qn.2] Why did you reach out to an AI chatbot here, rather than another resource (online or otherwise)?* [free-text response]

[Qn.3] How do you feel about how the AI chatbot responded?*

- Very positive
- Somewhat positive
- Neutral
- Somewhat negative
- Very negative

[Qn.4] Please explain why you feel this way. [free-text response]

[Qn.5] Do you have access to the transcript of this conversation and are willing to share it?*

↪ if No: (no follow-up)

↪ if Yes:

[Qn.5.1] What AI chatbot did you use? [free-text response]

[Qn.5.2] Approximately when did this conversation occur? [date: DD/MM/YYYY]

[Qn.5.3] Paste the transcript below. [free-text response]

885 $\hookrightarrow$ if No:

886 [Qn.4] Would you consider turning to an AI chatbot *[same topic as the stem]*?*

887 $\hookrightarrow$ if No:

888 [Qn.4.1] Why would you not consider reaching out to an AI chatbot here?*[free-text response]

890 $\hookrightarrow$ if Yes:

891 [Qn.4.2] Suppose you reached out to an AI chatbot *[same topic]*. What is a realistic question you might
892 pose?*[free-text response]

893 [Qn.4.3] Why would you consider reaching out to an AI chatbot here, rather than another resource?*[free-text
894 response]

895
896
897 *Note.* The label [Qn.4] is reused: in the Yes branch it is the free-text explanation; in the No branch it is the hypothetical-consideration
898 item. This mirrors the deployed instrument.
899

900 C.5 Section 1 — Mental Health

901 Each item applies the template of Appendix C.4 with topic:

902 [Q1] ...for advice on dealing with feelings of isolation, anxiety, depression, panic, or other mental health issues at work,
903 school, or another social setting.

904 [Q2] ...for advice on coping with a major life transition or a shift in how you see yourself.

905 [Q3] ...to help you figure out if you should start seeing a therapist for ongoing feelings of sadness or disconnection.

906 [Q4] ...for advice on navigating a situation where your personal values conflict with the expectations of your family or
907 community.

908 [Q5] ...for strategies to improve your body image or feel more comfortable in your own skin.

909 C.6 Section 2 — Gender & Sexual Health

910 [Q6] ...for help understanding your gender or sexual orientation.

911 [Q7] ...for help understanding the symptoms, causes, treatments, or other information about a sexually transmitted
912 infection (STI).

913 [Q8] ...for help understanding gender-affirming care.

914 [Q9] ...for information regarding the side effects or lifestyle impacts of taking gender or sexual health medications, such
915 as pre-exposure prophylaxis (PrEP) or estrogen.

916 [Q10] ...for tips on finding healthcare providers knowledgeable about your specific gender or sexual health needs.

917 C.7 Section 3 — Relationship Advice

918 [Q11] ...for advice on navigating communication, intimacy, or setting boundaries within a new relationship.

919 [Q12] ...for advice on exploring a new relationship dynamic that differs from your past experiences.

920 [Q13] ...for advice on communicating your needs to a partner who may not understand your perspective.

921 [Q14] ...for recommendations on books, movies, or media featuring representations of relationships like yours.

922 [Q15] ...for tips on physically or emotionally preparing for a new sexual experience.

C.8 Transcript “Donation” Section

Participants are asked whether they will upload their entire chat history for the AI chatbot they use most regularly, for a \$5.00 bonus.

↪ *if Yes:* After completing the survey, participants receive a Prolific message with a follow-up upload link (open $\approx$ 1 week). Stated validity requirements and data handling:

- First transcript message dated $\geq$ 4 weeks before the survey.
- $\geq$ 100 messages dated before the survey.
- One chatbot’s transcript only.
- Anonymized (proper names of people/places, addresses, phone numbers removed); never released publicly; shareable only with IRB-approved researchers under secure-handling plans.

Provider download tutorials referenced: ChatGPT/OpenAI, Claude/Anthropic, Gemini, Grok, Character.ai, Snapchat My AI, and Meta AI (Settings > Data & privacy > Manage your information > Download your information; “Meta AI app” and “All time” selected).

↪ *if No:* (no upload)

End-of-survey message. Return to Prolific and enter the completion code [Prolific Completion Code] to register the submission.

C.9 Follow-up Transcript Upload Survey

Re-asks willingness to upload (additional \$5.00), restating the validity and data-handling terms of Appendix C.8.

↪ *if Yes:*

[Q] Upload data as a single .zip file. [.zip upload]

[Q] Which AI chatbot did you upload a transcript for?

- ChatGPT/OpenAI
- Claude
- Gemini
- Grok
- Character.ai
- Snapchat My AI
- Meta AI
- Other (please specify) [free-text response]

End-of-survey message. A valid uploaded transcript is compensated \$5 via Prolific’s Bonus Payment feature.